

Frank Fu


UGC & TREND ANALYZER

Sentiment, keywords and themes
pulled from 50K+ Amazon reviews.


WHAT HAPPENS BETWEEN THE REVIEW AND THE ANSWER

Nothing here is a prompt to a chatbot. Four models do four narrow jobs, each one checkable on its own, and every judgement call is frozen to a file that lives in version control.

STAGE ONE — BUILD A TAXONOMY FROM THE DATA, NOT FROM A WORKSHOP



STAGE TWO — MEASURE EVERY SENTENCE AGAINST THAT FIXED TAXONOMY



THE SCORE TELLS YOU WHO LEFT. THE TEXT TELLS YOU WHY.

Four findings from this corpus that a rating-based read cannot produce — each with the decision it forces — then what the same method is for elsewhere.

17.5%

Complaints hide in five-star reviews
Of sentences inside 5★ reviews, 17.5% are complaints — 26.5% in 4★. Problems reported by customers who **stayed**.
SO DO THIS
Re-rank the backlog on sentence-level complaints rather than review scores. What your happiest customers complain about is the cheapest retention work available — they have not left yet.

+0.22★

Some complaints are demand signals
Complaining you cannot **find** a product travels with a rating 0.22★ **higher** (p=0.03). These come from devotees.
SO DO THIS
Route them to supply and channel, never to product. A rise in “can’t find it” is a restock trigger and a demand forecast — and it belongs on a different dashboard from defects.

84.6%

The worst category is a logistics problem
Protein snacks rate lowest (3.88★). Condition-on-arrival mentions are 84.6% negative, at 14% of reviews — dates, bags, descriptions.
SO DO THIS
Audit date rotation, seal integrity and listing accuracy before touching the formulation. A volume-ranked read would have sent a quarter of R&D spend at the recipe instead.

“natural”

A marketing claim that backfires
The strongest marker of a **low**-rated drink review — reviewers reach for it while giving one star. In pet food the equivalent is “china”.
SO DO THIS
Substantiate the claim or drop it from listing copy, and lead with the praise vocabulary instead. Check origin disclosure in pet lines before the next packaging run.

THE SAME METHOD, WHEREVER TEXT SITS BESIDE A SCORE

Live comment streams	Comments arrive faster than anyone can read them. Aspect plus polarity per comment turns a feed into a live count of which specific thing people are reacting to — a pricing backlash separated from a delivery outage within minutes, not at the next review cycle.	Silent mislabelling	A classifier that never abstains files chatter under a real theme. Guard: an explicit unassigned class — 20.7% here — and a hand-judged accuracy published beside every result.
Support & CSAT	Ticket text inherits the survey score, so friction inside “satisfied” tickets is never counted. Same correction; weight against renewal instead of the survey.	False causation	Polarity is read from the same text that produced the score, so the link is mechanical. A triage order, never “fix X and the rating rises”.
App store & marketplace	A five-star review naming a missing feature is the best product input a team receives, and exactly what a star-weighted dashboard discards.	Thin slices	Small segments produce confident nonsense. Guard: floors that make the system decline to answer — four of fourteen categories report no language analysis at all.
Surveys & complaints	eNPS, patient experience, regulated narratives. Aspects rather than moods means “discharge communication” can be tracked as improving or worsening over time.	Drift & fluent prose	Renamed themes destroy trends; generated summaries read well whether or not they are right. Guard: taxonomy frozen in git, and every number machine-checked against its evidence.

To point it at something else you need free text, a numeric outcome and a few thousand rows. Only the taxonomy is rewritten — an afternoon — and every stage after it is unchanged. Open weights, one laptop GPU, no API key, no per-call cost.