

Frank Fu

UGC & TREND ANALYZER


Sentiment, keywords and themes
pulled from 50K+ Amazon reviews.

23;6,GJ.Z//AEM:3;7SFR9FU6EGSBADZ#RT0HG3LMMIY0A;B2
XI4ILT,M.MY:DCQZ:7E4,RNE2CXIIV7IBWC.EUE8.IV40B#YA
BH,5AJMV2PQ1MYUICRHR07U,ERCRT#IMRJ.G/XM4A8LW7KOY
TUBFVRBECVI7725F2S979Y,J758ABCZ#FE3Z4V4G,QBJSW5;
.LUMMIQ10I.B#X6VM4C,#MNKGI3G3W:.M3XMGAF.#6VQHKZ
Z86/H#,PZMAQV67G/:CKK8D57TSRS SP.YUDU04PG3PH0NWHI
#MPP/CA1C8FK.4N5E2,NQJ/:9LX3F153AN:TUJUH2RT07XBZ
/188FGF//073VV5ECWY3.XH0P.KYSDHJL34CMN::2RT9QZVFC
U7,8#AJF,G3;6MF285TZTEIZHA4#ANJYTZFMHR.YJILZ3FCHI
QV2YVGK89,JM7MBMI#VXHI#9QVHAX#K93LIY.B0JHLI#DCJ54
D,YTE2Z478ZEX0K0EI83PBE2Z;TU54MDWB9.DTURI2KTMP SX
Q#SMRBJC;5E/UC6.78K12QP8;Q:DL:56QBL8PDC,Q56Q5:D3
C,QPFX,K/V;LZNDW94NT5LN0RR7Q#NXXPHU2Y7TPM:N1I92J
ENWX.15YSRZKPVJQ9BHER5TU EWXQLYXN.80E2YIRXBMQP8A5
08/U6EGT24,4QZ.UNI0IRI,0AIR.YIHM8HB4E/Z0#NJGQ1UR
009D,,MP2DB8EBT3NX7P58G40,#UXRX2R/GNNINGX1ZZ8NL;
2.#H33YY#C2;DJV,4ZI5L/JGQCQW81RP95ED9EJLCOH,,5/.
,TEX0E1,3HW0;HSF3402XSRA PXQR;GNXBWBJX6/OZOHE//7N
08G6EI6SP0NK3G6#0HAHY7ZT8CTQ059#GT:X4/2TDYZGC475
F;1XSG0VJSD5,I.,NA4.EFI2HEW9H/NUHRCY5Y4KUW/;Y;WJ
5LG6CW95WGJM,K;XXSQW8S95BRJK8,K;Q1X82W.#09EA#MTJ
9TOJJJB;#,K7CTBJ;:FENOVZP;VWP0J2V/:QB06NPF6N9XEY0
UIUMP3HJEWHLR827FO T6VG::5M4B;DCNB5ELKC,4UJ;135,K
Z##F./B2RPZQ.1D5J6M2IRL4PBV/Z/KST,Q,D0A55WKO,JUL
87MXUVW.NPF1096Q07LDLR;U/D14QDA,I5P/,FI#BZ,P5U.T



READ EVERY SENTENCE, NOT EVERY RATING

Most feedback systems attach a review's star rating to all of its text. That single shortcut is where the useful signal gets destroyed — so the pipeline scores each sentence independently, then rebuilds the picture from the bottom up.



THREE APPROACHES WERE TRIED AND DISCARDED BEFORE THIS ONE

REJECTED · DENSITY CLUSTERING

HDBSCAN returned 2 clusters and left 90% of sentences as noise. Keeping the conventional choice would have meant discarding nine tenths of the evidence.

REJECTED · CLUSTERS AS LABELS

56.2% accurate. Clustering must place every sentence somewhere, so unrelated chatter lands inside a theme and stays there.

REJECTED · DEFINITIONS AS PROMPTS

62.5% left unassigned. Long compound hypotheses entail badly; short label phrases fixed it with no change of model or data.

The result of doing it this way: 17.5% of sentences inside five-star reviews turn out to be complaints — criticism from customers who stayed, which a rating-based read cannot see at all. Two earlier versions of the assignment step scored 56.2% and 82.4%; both failures are documented rather than dropped.

ANY TEXT WITH A SCORE ATTACHED

Nothing here is specific to food, to Amazon, or to stars. Wherever an organisation collects free-text feedback next to a numeric outcome, the same shortcut is being taken — and the same buried signal is waiting.

SaaS support & CSAT

Friction described inside "satisfied" tickets is never counted. Aspect extraction surfaces it and ranks it against renewal, not against the survey score.

Patient experience

Complaints buried in high-scoring surveys — discharge instructions, waiting communication — are operational and cheap, precisely because the patient was satisfied.

HOW TO TELL WHETHER YOUR OWN DATA HAS THIS PROBLEM

TEST ONE

Pull fifty pieces of feedback with the **highest** scores and read them properly. Count how many name a specific problem. If the answer is not near zero, your dashboard is not showing it to you.

App store reviews

A five-star review naming a missing feature is the best product input a team gets, and exactly what a star-weighted dashboard discards.

Financial complaints

Regulated narratives need triage, not sentiment. A published error rate and an explicit "unassigned" class is what makes such a model auditable.

TEST TWO

Ask what your top theme would be if you ranked by **impact on the outcome** rather than by mention count. If nobody can answer, you are prioritising by volume.

Employee engagement

Comments paired with eNPS, same inheritance bug. Aspect themes let "manager support" be tracked as improving or worsening — which "sentiment down 4%" never tells you.

Hospitality & travel

"Loved the hotel, the air conditioning was broken" — one review, two facts. Multi-label assignment keeps both instead of averaging them away.

TEST THREE

Ask what your classifier's **accuracy** is, on a sample somebody actually checked. An unanswered question here means every number downstream is unmeasured.

What a new domain needs: free text, a numeric outcome, and a few thousand responses. Only the taxonomy is rewritten — an afternoon's work; every stage after it is unchanged. It runs on a laptop GPU with no API cost.
What it will not give you: causation — polarity is read from the same text that produced the score, so the ranking is a triage order, not a treatment effect.