

Frank Fu


UGC & TREND ANALYZER

Sentiment, keywords and themes
pulled from 50K+ Amazon reviews.


WHAT HAPPENS BETWEEN THE REVIEW AND THE ANSWER

Nothing here is a prompt to a chatbot. Four models do four narrow jobs, each one checkable on its own, and every judgement call is frozen to a file that lives in version control.

STAGE ONE — BUILD A TAXONOMY FROM THE DATA, NOT FROM A WORKSHOP




STAGE TWO — MEASURE EVERY SENTENCE AGAINST THAT FIXED TAXONOMY




EIGHT CATEGORIES, AND THE WORD EACH ONE IS LOSING ON

Bars are weighted log-odds — how decisively a word separates a one-star review from a five-star one, on a shared scale across all eight. Every word was read back against real low-rated sentences before it was allowed on this page.


■ drives complaints ■ drives loyalty Longer bar = harder evidence · shared scale, max z = 4.2 (“china”)




All five are **format or strength** — pods that don't fit the machine, packets that brew thin. An unset expectation, not bad coffee.




Two named failures: peach “tastes like soap”, and green tea bags hold **half the leaf** of rivals. Fill weight and two SKUs.




“Sea salt” is a **labelling dispute** — buyers say the flavour is named backwards. Oil and grease do the rest.




“Pieces” = **arrived broken**. “Diabetic” = sugar-free promised, not delivered. Neither failure is about flavour.




Watered down, strange aftertaste, **not worth the cost** — and “received” flags batch-to-batch inconsistency, not a recipe fault.



Buyers read the panel: **“first ingredient is mineral oil — olive oil is fourth.”** The front of pack and the back disagree.



The strongest signal in the entire dataset. Owners link country of origin to illness and stop buying. Sourcing is the product.



Lowest-rated category, and **no complaint word is about the recipe**. Stock arriving at its expiry date, bags split at the seam.

Same board, any channel Two models per message — which aspect, how negative — so the same board redraws for **live comments**, support tickets, app reviews or survey text.

Abstains when thin · 88.2% audited · no causal claims