

Frank Fu


UGC & TREND ANALYZER

Sentiment, keywords and themes
pulled from 50K+ Amazon reviews.


WHAT HAPPENS BETWEEN THE REVIEW AND THE ANSWER

Nothing here is a prompt to a chatbot. Four models do four narrow jobs, each one checkable on its own, and every judgement call is frozen to a file that lives in version control.

STAGE ONE — BUILD A TAXONOMY FROM THE DATA, NOT FROM A WORKSHOP



STAGE TWO — MEASURE EVERY SENTENCE AGAINST THAT FIXED TAXONOMY



THE SCORE TELLS YOU WHO LEFT. THE TEXT TELLS YOU WHY.

Every organisation with a rating already knows its number went down. None of them learn the reason from the number. This reads the sentence instead — which is where the reason always was.

Where the same problem exists today

Live comment streams	Comments arrive faster than anyone can read them. Aspect + polarity per comment turns a feed into a live count of which specific thing people are reacting to — separating a pricing backlash from a delivery outage within minutes, not at the next review cycle.
Support & CSAT	Ticket text inherits the survey score, so friction inside "satisfied" tickets is never counted. The same fix applies, and the outcome you weight against becomes renewal.
App store & marketplace	A five-star review naming a missing feature is the best product input a team receives, and exactly what a star-weighted dashboard throws away.
Employee & patient surveys	Free text beside eNPS or an overall score. Because themes are aspects rather than moods, "manager support" or "discharge communication" can be tracked as improving or worsening over time.
Regulated complaints	Narratives needing triage and an audit trail. A published error rate and an explicit "unassigned" class are what make such a model defensible to a reviewer.

What has to be true before anyone trusts it

Silent mislabelling	A classifier that never abstains will file chatter under a real theme. Guard: an explicit unassigned class, and a hand-judged accuracy figure published beside every result.
False causation	Polarity is read from the same text that produced the score, so the link is mechanical. This is a triage order, never "fix X and the rating rises". Stated in those words.
Thin slices	Small segments produce confident nonsense. Guard: sample-size floors that make the system decline to answer — four of fourteen categories here report no language analysis at all.
Drifting definitions	If a model renames themes on each run, no trend is real. Guard: the taxonomy is frozen in version control, so any change to a theme appears in a diff.
Confident prose	Generated summaries are fluent whether or not they are right. Guard: every numeric claim is machine-checked against the evidence, and the checker's own blind spot is documented.

What it takes to point this at something else: free text, a numeric outcome, and a few thousand rows. Only the taxonomy is rewritten — an afternoon — and every stage after it is unchanged. **The finding that makes it worth doing:** 17.5% of sentences inside five-star reviews are complaints. Those are the problems your happiest customers are telling you about, and the score is hiding all of them.