

Frank Fu


UGC & TREND ANALYZER

Sentiment, keywords and themes
pulled from 50K+ Amazon reviews.


WHAT HAPPENS BETWEEN THE REVIEW AND THE ANSWER

Nothing here is a prompt to a chatbot. Four models do four narrow jobs, each one checkable on its own, and every judgement call is frozen to a file that lives in version control.

STAGE ONE — BUILD A TAXONOMY FROM THE DATA, NOT FROM A WORKSHOP



STAGE TWO — MEASURE EVERY SENTENCE AGAINST THAT FIXED TAXONOMY



Audited, not asserted 102 sentences judged by hand: 88.2% theme accuracy (CI 80.6–93.1), 95.1% polarity. Two earlier designs scored 56.2% and 82.4% and are documented.	Claims checked mechanically Every number in every generated brief is parsed and matched to its evidence before it renders. 103 claims, 0 unmatched — and it has caught real errors.	Runs anywhere No API key, no vendor, no per-call cost. Every model is open-weight and runs on a laptop GPU; outputs are committed so the dashboard opens instantly.
---	---	---

THE SCORE TELLS YOU WHO LEFT. THE TEXT TELLS YOU WHY.

Four findings from this corpus that a rating-based read cannot produce — then what the same method is for.

17.5%

Complaints hide in five-star reviews

Of all sentences inside 5★ reviews, 17.5% are complaints — 26.5% in 4★. These are problems reported by customers who **stayed**, and a rating-weighted read counts every one of them as praise.

+0.22★

Some complaints are demand signals

Complaining you cannot **find** a product travels with a rating 0.22★ **higher** (p=0.03). Scarcity complaints come from devotees. Route them to distribution, not to the fix list.

84.6%

The worst category is a logistics problem

Protein snacks rate lowest (3.88★). Its condition-on-arrival mentions are 84.6% negative, at 14% of reviews — complaints about dates, bags and descriptions. Nothing in the recipe needed changing.

“natural”

A marketing claim that backfires

The single strongest marker of a **low**-rated drink review. Reviewers reach for the word while giving one star — the claim was read as a promise about taste and missed it. In pet food the equivalent word is “china”.

THE SAME METHOD, WHEREVER TEXT SITS BESIDE A SCORE

Live comment streams	Comments arrive faster than anyone can read them. Aspect plus polarity per comment turns a feed into a live count of which specific thing people are reacting to — a pricing backlash separated from a delivery outage within minutes, not at the next review cycle.	Silent mislabelling	A classifier that never abstains files chatter under a real theme. Guard: an explicit unassigned class — 20.7% here — and a hand-judged accuracy published beside every result.
Support & CSAT	Ticket text inherits the survey score, so friction inside “satisfied” tickets is never counted. Same correction; weight against renewal instead of the survey.	False causation	Polarity is read from the same text that produced the score, so the link is mechanical. A triage order, never “fix X and the rating rises”.
App store & marketplace	A five-star review naming a missing feature is the best product input a team receives, and exactly what a star-weighted dashboard discards.	Thin slices	Small segments produce confident nonsense. Guard: floors that make the system decline to answer — four of fourteen categories report no language analysis at all.
Surveys & complaints	eNPS, patient experience, regulated narratives. Aspects rather than moods means “discharge communication” can be tracked as improving or worsening over time.	Drift & fluent prose	Renamed themes destroy trends; generated summaries read well whether or not they are right. Guard: taxonomy frozen in git, and every number machine-checked against its evidence.

To point it at something else you need free text, a numeric outcome, and a few thousand rows. Only the taxonomy is rewritten — an afternoon's work — and every stage after it is unchanged: same embeddings, same entailment, same regression, same claim checker. It runs on open weights, on one laptop GPU, with no API key and no per-call cost.