

Frank Fu


UGC & TREND ANALYZER

Sentiment, keywords and themes
pulled from 50K+ Amazon reviews.


WHAT HAPPENS BETWEEN THE REVIEW AND THE ANSWER

Nothing here is a prompt to a chatbot. Four models do four narrow jobs, each one checkable on its own, and every judgement call is frozen to a file that lives in version control.

STAGE ONE — BUILD A TAXONOMY FROM THE DATA, NOT FROM A WORKSHOP



STAGE TWO — MEASURE EVERY SENTENCE AGAINST THAT FIXED TAXONOMY



Audited, not asserted 102 sentences judged by hand: 88.2% theme accuracy (CI 80.6–93.1), 95.1% polarity. Two earlier designs scored 56.2% and 82.4% and are documented.	Claims checked mechanically Every number in every generated brief is parsed and matched to its evidence before it renders. 103 claims, 0 unmatched — and it has caught real errors.	Runs anywhere No API key, no vendor, no per-call cost. Every model is open-weight and runs on a laptop GPU; outputs are committed so the dashboard opens instantly.
---	---	---

WHAT CUSTOMERS CARE ABOUT, CATEGORY BY CATEGORY

Estimated drag on the star rating — how often a complaint occurs here × what it is associated with costing. Darkest cell in a column is where that lever is worth pulling.

	rev	Taste	Ingredients	Condition	Price	Packaging
coffee	634	0.16	0.03	0.05	0.02	0.02
snack	558	0.14	0.04	0.05	0.02	0.02
drink	436	0.15	0.06	0.04	0.03	0.02
baby & pet	431	0.11	0.03	0.03	0.02	0.02
tea	334	0.16	0.03	0.02	0.02	0.02
candy	285	0.18	0.03	0.05	0.02	0.02
staples	256	0.14	0.05	0.04	0.02	0.02
other	249	0.08	0.02	0.06	0.03	0.02
condiments	236	0.20	0.07	0.05	0.03	0.02
protein snacks	137	0.17	0.04	0.12	0.04	0.02
seasoning & spice	112	0.17	0.07	0.04	0.01	0.01
chocolate	103	0.18	0.04	0.07	0.02	0.01
baked goods	91	0.15	0.05	0.12	0.02	0.02
dairy	74	0.12	0.07	0.05	0.01	0.01

Figures are stars of drag per review. **Shading runs within each column**, so every lever shows its own worst category — taste would otherwise flatten the rest. Dotted = under 15 reviews, read with caution. · = below the floor, not estimated.

Taste is table stakes, not a strategy

-0.12 to -0.20 in every category

It tops all 14, so it cannot prioritise anything. The decisions live in the columns beside it — which is exactly what a share-of-voice ranking hides.

Protein snacks: fix the warehouse

Condition -0.117 · n=19 · 84.6% negative

Three times the condition drag of coffee. Dates, seals, listing accuracy — not the recipe. Audit fulfilment before R&D spends a quarter.

Condiments & drinks: the label, not the liquid

Ingredients -0.071 and -0.057 · n=51, n=76

Highest ingredient drag anywhere, on solid samples. Buyers read the panel and feel misled. Reformulate or restate — “natural” is the top complaint word in drinks.

Availability complaints are upside

+0.22★, p=0.03

The only theme that moves ratings **up**. Restock signal, not a defect — route it to supply, keep it off the fix list.

Live comments

Same two models per comment: which thing is being reacted to, right now.

Support tickets

Friction inside “satisfied” tickets, weighted against renewal.

App reviews

Missing features named in 5★ reviews.

Surveys

eNPS and patient text, tracked per aspect over time.

Guardrails

Abstains when thin, publishes its error rate, claims no causation.